

ChCore Lab 1: Booting a machine

517021910851-于亚杰

Part A: Bootstrap

Exercise 1

Read or at least carefully scan the Part A1 and A3 of the Arm® Instruction Set Reference Guide to be familiar with ARM ISA. Part D is a useful reference for ARM instructions. You can scan the section 1 and 2 in this guide to be familiar with gcc AArch64 grammar quickly. Please write down some differences compared with x86-64 you have learned before.

Some differences I've observed are:

- ARM uses a RISC architecture while X86 uses a CISC architecture
- load and store instructions are incorporated in X86
- X86 has smaller code sizes
- instruction length of X86 is variable while it's fixed in ARM
- ARM heavy use of ram
- X86-64 has fewer registers
- A “word” on x86 is 16-bits and a “doubleword” is 32-bits. A “word” for ARM is 32-bits and a “doubleword” is 64-bits.
- no branch prediction(ARM) vs. branch prediction(X86)

Exercise 2

Use GDB's where command to trace the entry (first function) and its address of the boot loader. Where is this function defined? You can also use readelf -S build/kernel.img to read the symbol table.

```
(gdb) where
#0  0x0000000000008000 in _start ()
-----
> readelf -S build/kernel.img
[ 1] init                PROGBITS          0000000000008000 00010000
      0000000000049680 0000000000000008 WAX          0      0      4096
```

So the function is defined in the init section of the ELF and its starting address is 0x80000.

PartB: The Boot Loader

Exercise 3

Before we introduce multi-processors, we focus on a single processor. In Raspi3, all processors boot simultaneously. Please read the boot code in boot/start.S and explain how to hang other processors according to the control flow.

```
mrs    x8, mpidr_el1
```

The code first get the processor id, and store it in register x8.

```
mov     x9, #0xc1000000
bic     x8, x8, x9
cbz     x8, primary
```

Then use conditional branch to differentiate one processor from secondary processors.

If $x8 \text{ AND } \text{not } 0xc1000000 == 0$, the code will switch to Primary part to continue booting which happens to only one processor. And other processors will go on to secondary_hang part and just hang there forever.

Exercise 4

You have learned basic pointer operations in the ICS course. Then download the code for pointer-warmup.c, compile by `gcc pointer-warmup.c`, run `./a.out`, and make sure you understand all of the outputs. How all the values in printed lines 2 to 4 get there, and why the values printed in line 5 are seemingly corrupted.

line 1: print address of the pointers and each variable has been allocated 8 bytes space

line 2: `a[i]` is assigned long value

line 3: `*(c + 2)` refers to `c[2]` and `3[c]` refers to `c[3]`

line 4: `c = c + 1`; means the address of `c` is increased by 8 bytes

line 5: `c = (long *) ((char *) c + 1)`; means the address of `c` is increased by 1 bytes and `*c` still is a long value occupying 8 bytes. And change the value of `*c` will change the first byte of `a[2]`. By the way big endian storage should also be mentioned here.

line 6: `b = (long *) a + 1`; address of `b` = address of `a` + 8 bytes; `c = (long *) ((char *) a + 1)`; address of `c` = address of `a` + 1 byte;

Exercise 5

Read the objdump information of build/kernel.img, the code of boot loader is stored in .init section, and codes of the kernel are stored in .text, explain the differences of their LMA and VMA. And try to explain why these differences happened.

`.init` section has the same LMA and VMA while `.text` section has different ones. Since the boot_loader is started first, there's no need to extinguish the two addresses and it's faster to directly boot from the load address. And more importantly the kernel is run later so in order to save the RAM the code of the kernel can stay in ROM for a while and copied to the RAM when it's needed in run time. Another reason may be that before the kernel space is setup the code has to loaded in the user space, and then copied to kernel space.

Part C: The Kernel

Exercise 6

`tfp_printf` and `printk` are the interface function, and they support format print similar to standard `printf`. While we omit the code to print numbers in different bases (e.g., 2, 8, 10, 16), please fill the functions `write_num` in `inboot/print/printf.c` and `printk_write_num` in `kernel/common/printk.c` to implement this function. They have the same logic.

divide `n` by `base` until it comes to 0

store the remainder from the end of the array to the beginning, if `neg` is not zero, print '-' first

then call `write_string` with argument `str` referring to the position of the most significant digit of `n`

Exercise 7

Where do codes of kernel stack initialization lay? Which function does the kernel initialize its stack? Where is its stack located in the memory? How does the kernel reserve space for its stack?

```
(gdb) x/g $sp
0xffffffff0000d20f0 <kernel_stack+8192>: 0x0000000000000000
(gdb) p &kernel_stack
$1 = (char (*)[4][8192]) 0xffffffff0000d00f0 <kernel_stack>
```

The codes of kernel stack initialization lay in `kernel/head.S`. function `start_kernel` will initialize its stack. The stack size is initialized to 8KB(`KERNEL_STACK_SIZE`), and the stack bottom address is `0xffffffff0000d20f0` while the stack top address is `0xffffffff0000d00f0`.

```
ldr    x2, =kernel_stack
add    x2, x2, KERNEL_STACK_SIZE
mov    sp, x2
```

kernel_stack+KERNEL_STACK_SIZE is assigned to the stack pointer and thus the space is reserved.

Exercise 8

To become familiar with the C calling conventions on the AArch64, find the address of the test_backtrace function in kernel/main.c by gdb, set a breakpoint there, and examine what happens each time it gets called after the kernel starts. How many 64-bit words does each recursive nesting level of test_backtrace push on the stack, and what are those words?

```
(gdb) p &test_backtrace
$1 = (void (*)(long)) 0xffffffff00000cc020 <test_backtrace>
```

The function address of test_backtrace is 0xffffffff00000cc020 .

```
(gdb) x/10g $sp
0xffffffff00000d2060 <kernel_stack+8048>: 0xffffffff00000d2080      0xffffffff00000cc070
0xffffffff00000d2070 <kernel_stack+8064>: 0x0000000000000003      0xffffffffffffffff
```

Every time test_backtrace calls itself the stack is decreased by 32Bytes. So that is four 64-bit words pushed on the stack. And the four words in sequence are frame pointer(last function), return address(the address to return to), argument x and 0xffffffffffffffff (this value is changed by function printk_format and are not used in function test_backtrace).

Exercise 9

In AArch64, return address (x30), frame pointer (x29), and arguments are transmitted by registers. However, when a caller function calls a callee function, these registers should be stored in the stack, where do these values store?

At the beginning of the callee function, a stp instruction is executed to push the frame pointer(x29) and return address(x30) onto the stack. As is shown in the following code:

```
# x0 is the argument passed here
# push x30 x29
0xffffffff00000cc020 <test_backtrace>      stp      x29, x30, [sp, #-32]!
# store a callee-saved register
0xffffffff00000cc024 <test_backtrace+4>    mov      x29, sp
0xffffffff00000cc028 <test_backtrace+8>    str      x19, [sp, #16]
```

To be more informative, x30 is pushed prior to x29. And the instruction above is not the common case since sp grows by 32 bytes while two registers combined are 16 bytes. So the instruction is

practically equivalent to: (1) $sp = sp - 16$ (2) push x30 (3) push x29.

The excessive bytes are left for function arguments.

change the function to accept 5 arguments and they are shown to be pushed in the following way

(gdb) x/10xg \$sp

0xffffffff0000d1fe0 <kernel_stack+7920>: 0xffffffff0000d2020	0xffffffff0000cc0a0
0xffffffff0000d1ff0 <kernel_stack+7936>: 0x0000000000000003	0x0000000000000011
0xffffffff0000d2000 <kernel_stack+7952>: 0x0000000000000012	0x0000000000000013
0xffffffff0000d2010 <kernel_stack+7968>: 0x0000000000000014	0xfffffffffffffffff

The function arguments are pushed onto the stack(in reversed order) before the callee function begins. By the way, sp is 16 bytes-aligned so sometimes there could be a 8-bytes space waste.

Exercise 10

Implement the backtrace function `mon_monitor` in `kernel/monitor.c` as specified above. Use the same format as in the example; otherwise, the grading script will be confused. Because of compiler optimization, you only need to consider the case of `test_backtrace`.

As is shown in the last exercise, the stack from bottom to top is like:

args(if #args is odd, add 0xffffffff for alignment intention), fp0, lr0, args, fp1, lr1.....

So by the help of `read_fp()`, it's easy to locate all the frames because `*fp` would refers to the position of the previous frame and the five arguments.

The last problem is about when to stop tracing. From `boot/start.S`, I found that the stack pointer is set to `0x1000+boot_cpu_stack`, which is `0x88530`, before calling `init_c`. And also the frame address of the outer-most function is actually `0x0`. So I guess both ways are reasonable to stop tracing when `*fp` equals 0 or `fp + 16` equals `0x88530`.